

Processes 1

Jonathan Windle

University of East Anglia

J.Windle@uea.ac.uk

June 9, 2017

Overview I

1 POSIX

2 Multiprogramming

- Three Views
- Side Effects
- Process vs Program
- Process Creation
- Process Termination
- Process Hierarchies
- Process States
- Process Model
- Implementation of Processes
- Modelling Multiprogramming
- Execution Modes

3 Threads

- Single-Threaded vs Multithreaded Processes
- Advantages of Threads

Overview II

- Thread States
- Implementing Threads
- Threads in User Space
- Threads in Kernel Space


- For programs to run on any UNIX system IEEE developed a standard for UNIX called POSIX.
- Defines a minimal system cll interface comprising of about 100 procedure calls.
- Grouped into four categories:
 - Process management
 - File management
 - Directory and file system management
 - Miscellaneous.

Multiprogramming

- CPU switches from program to program running each for 10s of milliseconds.
- Creates an illusion of parallelism sometimes called pseudo-parallelism
- System designers have developed a conceptual model called sequential processes to keep track of multiple parallel activities.
- Most programs are sequential processes.
- They comprise of a series of instructions, executed sequentially.
- Assuming the input is unchanged they will always produce the same result.

Three views of Multi-programmings

- Computer multiprogramming four programs in memory.
- Four processes each with its own flow of control.
- Trace of execution shows all processes making progress.



- Because of CPU switching, the rate at which a process performs its computation will not be uniform and probably not even be reproducible if the same processes are run again.
- Processes must **NOT** be programmed with built in assumptions about timing.
- When a process has a particular critical real-time requirement meaning particular events must occur within a specified number of milliseconds, special measures must be taken.

Process vs Program

- Key idea is that a process is an activity of some kind.
- It has a:
 - Program
 - Input
 - State
 - Output
- A single processor may be shared amongst several processes.
- A process is a program in execution.
- A program is just a list of instructions in memory.

- OS need some way to create and terminate processes.
- Four principle events cause processes to be created:
 - System initialisation
 - Execution of a process creation system call (by another process)
 - A user types a command.
 - A user runs a batch job.
- In UNIX there is only one system call for creating a process: `fork`.
- `fork` creates an exact clone of the calling process. After `fork` there are two processes, the parent and the child, each having:
 - Their own **distinct address space**.
 - The same **memory image**.
 - The same **environment settings**
 - The same **open files**
- Usually the child process then executes `execve` to change its memory image and run a new program.


Process Termination

- Terminate due to:
 - Normal exit (voluntary).
 - Error exit (voluntary)
 - Fatal error (involuntary)
 - Killed by another process (involuntary).

Process Hierarchies

- Some systems, when a process creates another process, the parent process and child process continue to be associated.
- In UNIX, a process and all its children form a process group e.g.:
 - When a user sends a signal from the keyboard it is sent to all members of the process group currently associated with the keyboard.
 - When UNIX is started, `init` is present in the boot image. When it starts running, it reads a file telling how many terminals there are and forks one new process per terminal. It waits for a user to log in and if login is successful the login process executes a shell to accept commands. Hence, in UNIX all the processes in the whole system belong to a single tree, with `init` at the root.
- `ps tree` command shows the running processes as a tree.


Process States



1. Process blocks for input
2. Scheduler picks another process
3. Scheduler picks this process
4. Input becomes available

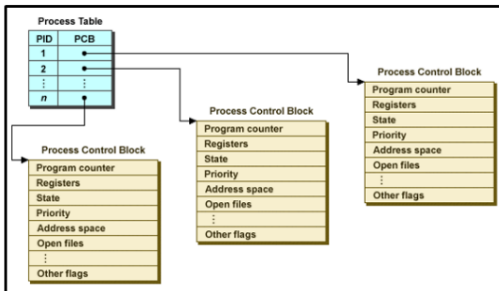
- Transition 1 occurs when the S discovers that the process cannot continue right now.
- Transition 2 & 3 are caused by the scheduler (part of OS).
- Transition 4 occurs when the external event for which the process was waiting (such as arrival of input) happens.
- External events send electrical signals called interrupts to the scheduler.

Process Model




- Using process model, becomes much easier to think about what's going on inside the system.
- In the process model, the lowest level of the OS is the **scheduler** with a variety of processes on top of it.
- All of the details of interrupt handling and actually starting and stopping the processes is hidden away in the **scheduler**.

Implementation of Processes



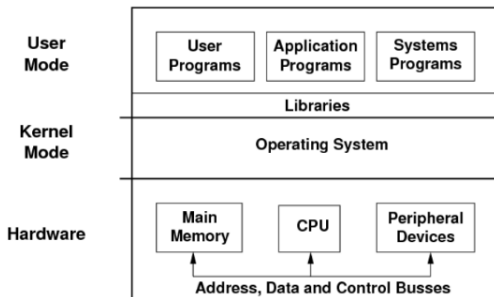
- The OS maintains a table called the **process table**.
- Individual entries in the table call **process control blocks** hold important information about the process state:
 - Program counter
 - Stack pointer
 - Memory allocation
 - Open files
 - Other accounting and scheduling information

Modelling Multiprogramming



- Model CPU usage from a probabilistic viewpoint.
- Suppose that a process spends a fraction of its time p waiting for I/O to complete.
- Assume n processes in memory at once.
- Probability that all n processes will be waiting for I/O is p^n ,
- Hence, CPU Utilisation = $1 - p^n$

Execution Modes



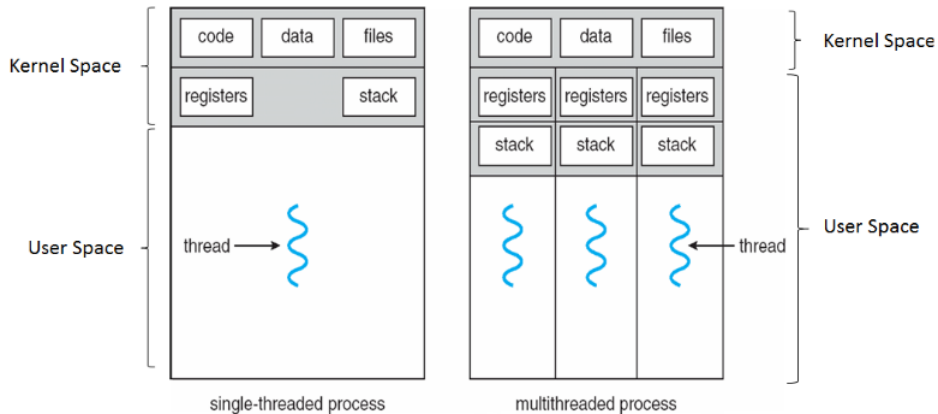
- Most computers have two mode of execution:
 - Kernel (or supervisor) Mode.
 - User Mode
- Processes running in kernel mode have unrestricted access to the machine:
 - They can disable all interrupts
 - Manipulate stacks
 - Unrestricted access to memory and I/O.

Threads

- In traditional OS, each process has an address space and a single thread of execution.
- It is sometimes desirable to have multiple threads of control in the same address space running in quasi-parallel.
- Threads run as though they were (almost) separate processes (except for the shared address space).
- A thread can be considered a **lightweight** process.

Single-Threaded vs Multithreaded Processes

Multithreaded partition of kernel / User space assumes threads are implemented in user space.



Advantages of Threads



- Decomposing the application into multiple threads that run in quasi-parallel the programming model becomes simpler. Also, threads share the same address space and data (unlike processes).
- Since threads are lighter weight than processes they are easier to create and destroy than processes. Creating a thread is typically 10-100 times quicker than creating a process.
- Threads are useful on systems with multiple CPUs.
- Threads yield no performance gain when all of them are CPU bound.
- Multiple threads are appropriate when they are part of the same job and actively cooperating with one another.

Per-process items	Per-thread items
Address space	Program counter
Global variables	Registers
Open files	Stack
Child processes	State
Pending alarms	
Signals and signal handlers	
Accounting information	

- A thread can be in one of three states:
 - Running
 - Blocked
 - Ready
- Different ways it can be blocked:
 - When the process blocks, ALL threads block
 - Threads may block waiting for an event
 - Threads may enter a sleep state to be woken by a signal from another thread.

Implementing Threads

Threads can be implemented in user space or kernel space



Threads in User Space

- Advantages:

- As far as the kernel is concerned it is running one single threaded process.
- Can be run on OS that doesn't support multithreading.
- Thread switching can be made very fast.
- Each process can have its own scheduling algorithm for threads.

- Disadvantages:

- If a thread starts running, no other thread in that process will ever run unless the first thread gives up the CPU.
- Blocking system calls are difficult to implement.
- Programmers generally want threads precisely in applications where the threads block often, and it's more efficient to implement as a process.

Threads in Kernel Space

- Advantages:
 - Kernel threads do not require any new, nonblocking system calls.
- Disadvantages:
 - The cost of creating and destroying threads in the kernel is substantial.
 - All calls that might block a thread are implemented as system calls at considerable greater cost than a call to a run-time system procedure.
- Other issues:
 - What happens when the process forks?
 - How can we handle signals?

Summary

- Processes are Ready,Running or Blocked.
- Threads may also Sleep and Wait.
- Two modes of execution User & Kernel processes switch to kernel mode during system calls.
- Threads are lightweight processes.
- Two types of thread implementations:
 - Kernel threads: Managed by the OS
 - User threads: Managed by user.

The End