

B Trees

Jonathan Windle

University of East Anglia

J.Windle@uea.ac.uk

June 3, 2017

1 Intro

2 Inserting


- Inserting - Case 1
- Inserting - Case 2
- Inserting - Case 3
- Inserting - Case 4

3 B-Trees

- Comparison

Intro


- Let K be a totally ordered set.
- A **2-3 tree**, t is a tree on $K \cup (K \times K)$ possessing the following properties:
 - Each non-leaf node has **2 or 3 children**
 - If p is a node with 2 children then it contains one key such that:
 - All keys in the left subtree of $p < \text{key}(p) < \text{all keys in the right subtree of } p$.
 - If p is a node with 3 children, then it contains two keys, $ltkey$ and $rtkey$, such that:
 - All keys in the left subtree of p :
 - $< ltkey$
 - $< \text{all keys in the middle subtree of } p$
 - $< rtkey$
 - $< \text{all keys in right subtree of } p$.
 - All leaf nodes are at the same level.



Inserting - Case 1

Insertion in a leaf node containing a single key:

E.g. Insert 25:




Inserting - Case 2

Node whose children should contain new key has only two children, both with two keys:

E.g. Insert 55:



- Go to right child of node containing 51; inserting 55 in this node would give 55,67,80;
- Make the middle key of these 3 keys the right key of the parent;
- Split the other two values into two nodes.



Inserting - Case 3

Node whose children should contain new key already has 3 children, and new key cannot be inserted in a leaf containing just 1 key.

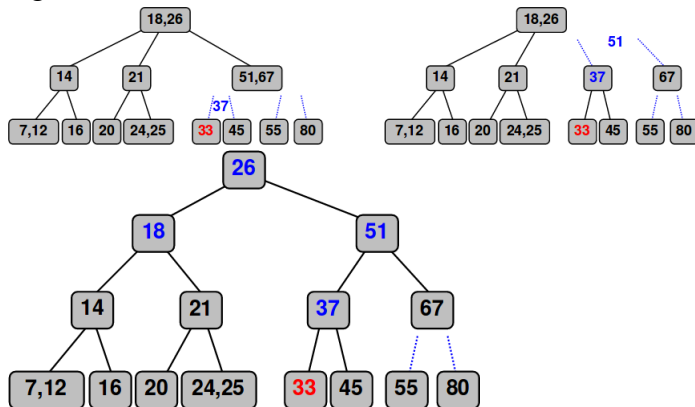
E.g. insert 16:



Inserting - Case 4

Splitting the root

E.g. Insert 16 in



- A 2-3 tree is a B-Tree of **order 3**.
- More generally, a **B-tree of order m** is an m -way search tree such that:
 - If the root is not a leaf node then it has between 2 and m children.
 - All non-leaf nodes except the root have between $\lceil m/2 \rceil$ and m children.
 - All leaf nodes are at the same level.
- The path length when searching in a B-tree is bounded above by:
 $\log_{\lceil m/2 \rceil} \left(\frac{n+1}{2} \right)$

Comparison of Height-Balanced Binary Search Trees and B-Trees

- The expected path length in a height-balanced BST containing n keys is:
$$\lceil \log_2(n+1) \rceil - 1$$
- Height of a B-tree of order m containing n keys is:
$$\leq \log_{\lceil m/2 \rceil} \left(\frac{n+1}{2} \right).$$
- It is easier to code BST algorithms
- B-Trees waste space if used for RAM-based dictionaries.
- B-Trees on disk units have $m = 256$ or more:
shallow tree \Rightarrow fewer disk accesses

The End