

Programming 2 Revision

Jonathan Windle

University of East Anglia

J.Windle@uea.ac.uk

May 26, 2017

1 Threads

- Parallel structures
- Processes & Threads
- Concurrent Programming
- Java Threads
 - overview
 - Declaring Threaded Objects
 - Creating and Starting threaded objects
 - run() vs start()
 - Issues surrounding writing threaded code
- Thread Interaction 1
 - join
 - Example of join()
 - Thread interaction
 - Naively parallel algorithms
 - Parallel sort
- Sleep/Interrupt
- Thread Interaction 2

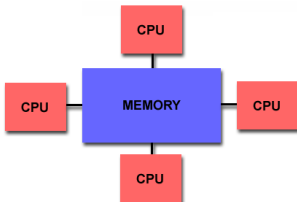
Overview II

- Synchronization
 - Claiming monitor lock
 - Volatile Variables
 - Monitor Locks vs Semaphores
- Thread Interaction 3
 - Wait Notify
 - wait()
 - notify() and notifyAll()
 - Wait/Notify-Textual Example
 - Wait/Notify-Code example
 - Wait/Notify Common Problems
 - Wait vs Sleep
- Common Problems with Threads
 - Deadlock
 - Starvation
 - Livelock
- Summary

Parallel Structures


Shared Memory

- Different processes access the same memory concurrently
- Each **process** shares memory with other **processes**.
- Structure adopted by multi core PC units now commonly available.



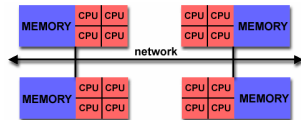
Distributed Memory

- Each CPU has its own memory.
- Messages passed between processes to coordinate their output.
- Structure adopted by High-Performance Computing (HPC) facilities.



Combined Shared/Distributed Memory

- Some CPUs share memory, some communicate via messages.
- Local processes also communicate with other groups via messaging.



Processes & Threads

- When program is launched, a **process** is created.
- **Processes** can be sent to different processors by message passing scheme.
- A **thread** is a **portion** of a **process** that can run **independently and concurrently** with other portions of the **process**.
- **Threads** within a **process** can share data.
- **Threads** rely on the OS (or other program) to determine what processor they are sent to.

Concurrent Programming

- Involves development of algorithms for running **simultaneously** in **multiple threads**.
- Java is inherently concurrent (Basic specification allows concurrent code development).
- C++ needs special library (e.g. pthreads) **C# is threaded**.
- Java does not allow for user controlled parallel processing.

Overview

- Any Java program has at least one thread, **main**.
- Any Java thread of execution is associated with an **instance** of the **Thread** class.
- Before new thread can start, a new instance of the **Thread** class must be **instantiated**.
- A java thread class implements the **Runnable** interface. Therefore every **Thread** instance has a method:

```
|| public void run() {...}
```

- When the **Thread** is **started** the code in the body of the **run()** method is executed.

Declaring Threaded Objects

Extends Thread

```
public class MyThread extends Thread{  
    public void run(){  
        for(int i = 0; i < 13; i++)  
            System.out.println  
                ("Thread Iteration " + i);  
    }  
}
```

- **Not** preferred, not defining the threads behaviour, just giving it something to run.
- Cannot inherit from this method.
- Each of your thread creates unique object and associate with it

Implements Runnable

```
public class MyRunnable implements Runnable{  
    public void run(){  
        for(int i = 0; i < 13; i++)  
            System.out.println  
                ("Runnable Iteration " + i);  
    }  
}
```

- Preferred method via composition with **Runnable**.
- Extendible, can have children classes.
- When you implement Runnable, it shares the same object to multiple threads.

Creating and Starting threaded objects

Starting Extends Thread

```
public static void main(){  
    MyThread x = new MyThread();  
    x.start();  
}
```

- Call **start**, not **run**

Starting Implements Runnable

```
public static void main(){  
    MyRunnable x = new MyRunnable();  
    (new Thread(x)).start();  
}
```

- Create the runnable object and pass it to a new Thread and call start on that.

run() vs start()

- Calling `run()` would just execute the code `sequentially`, essentially as a normal method.
- `start()` is a method defined in the `Thread` class and that calls `run()` method in separate `threads` that will run `concurrently` and thus thread the process.

Issues surrounding writing threaded code

- Developing the algorithms can be tough.
- Controlling the interaction of different threads:
 - Pausing for other threads to finish (**join**).
 - Pausing for a set time or until an exception to occurs (**sleep/interrupt**).
 - Waiting for another thread to notify that it's ok to continue (**wait/notify**).
- Controlling access to shared memory (**synchronize**).

Join()

Used to wait for another thread to finish

```
|| void join()
```

- This thread would wait for the calling thread to die before continuing.

```
|| void join(long millis)
```

- Waits at most, millis milliseconds for the thread to die.

```
|| void join(long millis, int nanos)
```

- Waits at most millis milliseconds plus nanos nanoseconds for the thread to die.

If two threads are waiting for the other to finish the process then it will never finish.

Example of join()

Code where the `void main()` method waits for two threads.

```
public static void main(string[] args) {
    MyThread t1 = new MyThread("Thread 1",1);
    MyThread t2 = new MyThread("Thread 2",2);
    t1.start();
    t2.start();
    try{
        t1.join();//Main now waits for t1 to die
        t2.join();//Main now waits for t2 to die
    }catch(InterruptedException ex){
        System.exit(0);
    }
    /*****
    * This for loop will only start when t1 and t2 have finished executing
    *****/
    for(int i = 0; i < 100000; i++) {
        if(i%1000==0)
            System.out.println("Main thread iteration: " + i);
    }
}
```

The `main` thread has been forced to wait for `t1` and `t2` to finish, but `t1` and `t2` are still running **concurrently** and therefore run completely independently.

Enforces **sequential** behaviour on **concurrent** code

Thread interaction

Threads can interact with each other either through:

- Storing references to each other (and calling methods on those references)
- Accessing shared memory: synchronisation.
 - Through the use of static variables within a threaded class.

```
public class MyThread {  
    public static ArrayList<Integer> sharedMem = new ArrayList<>();  
  
    public void run() {  
        for(int i = 0; i < 400; i++) {  
            if(!sharedMem.contains(i))  
                sharedMem.add(i);  
        }  
    }  
}
```

The static variable, `sharedMem` is accessible through any instance of `MyThread`, if multiple instances of `MyThread` exist and `start()` is called in quick succession, then all in this case will interact with `sharedMem` in a similar way.

Naively parallel algorithms

- If an algorithm can be split into sub-problems, then recombine the results from sub-problems, a thread system could be applied.
- If the threads do not interact with each other, then it is said that the problem is **Naively parallel**.

Concurrent sort(comparable[] ar, int n)

- 1 Split array into subarrays.
- 2 Create a thread to sort each subarray.
- 3 Merge sub arrays into sorted array.

Parallel sort

```
public class ThreadSort extends Thread {
    public static Comparable data[]; //Shared across all threads
    int start; //Local start position
    int end; //Local end position

    public ThreadSort(int s, int e) {
        start = s;
        end = e;
    }

    public void run() {
        Arrays.sort(data, start, end);
    }
}

public static void main(String [] args) {
    int m = 100;
    int n = 10;
    Comparable[] d = new Comparable[m];
    ThreadSort[] arr = new ThreadSort[n];
    for(int i = 0; i < n; i++) {
        arr[i] = new ThreadSort(i*m/n, (i+1)*m/n);
        arr[i].start();
    }
    //Main should join to ensure all sub sorts are complete before executing merge loop
    for(int i = 0; i < n-1; i++)
        merge(d, 0, (i+1)*m/n, (i+1)*m/n, (i+2)*m/n);
}
```


- `sleep()`
 - This pauses the thread for a fixed time or an `InterruptedException` is thrown when another thread has interrupted the current thread.
- `interrupt()`
 - Interrupts the thread called on. Uses an internal flag called the **interrupt status**. Calling the `interrupt` method sets the flag.
 - This only sets the internal flag, it doesn't stop the thread, or do anything if the thread is not sleeping.

Synchronization

- **Synchronization** involves locking an object so that it can only be accessed by one thread at a time.
- Allows enforcing **mutual exclusion** between threads.
- **Mutual exclusion** is implemented through **monitor locks**. All objects in java have a **monitor lock** that can be owned by other objects.
 - If one object owns the **monitor lock** of another object, then it has exclusive access to that object, until it gives up the lock.
- **Synchronization** is enforced using the **synchronized** keyword. Any **synchronized** method or statement can only be accessed by one thread at any time.
- When a thread invokes a **synchronized** method, it automatically acquired the **monitor lock** for that method's object and release it when the method returns.

Claiming monitor lock

- Direct **synchronized block**:

```
private SharedMemClass names;  
public void run() {  
    synchronized(names){ //Nothing else can touch names until after this block  
        names.alter();  
    }  
}
```

- **synchronized methods**:

```
public class SharedMemClass {  
    synchronized public void alter(){}//Synchronized method  
}  
public class Modifier extends Thread() {  
    private SharedMemClass names;  
    public void run() {  
//Instance of Modifier claims Monitor lock because of synchronized method  
        names.alter();  
    }  
}
```

block vs methods

Sometimes a method only needs mutual exclusion for a part, such as to read the memory, then update the memory. Only updating the memory needs to be **synchronized** via a **block**, opposed to the whole **method**.

Claiming monitor lock cont.

- Direct class monitor lock:

```
public class SharedMemClass {
    synchronized public void alter(){} //Synchronized method
    public static ArrayList<String> str = new ArrayList<>();
}
public class Modifier extends Thread() {
    private SharedMemClass names;
    public void run() {
        //This instance of Modifier claims lock on the whole class, locking all static variables
        //of the class
        synchronized(SharedMemClass.class);
    }
}
```

- Synchronized static methods:

```
public class SharedMemClass {
    synchronized static public void alter(){} //Synchronized method
    public static ArrayList<String> str = new ArrayList<>();
}
public class Modifier extends Thread() {
    private SharedMemClass names;
    public void run() {
        //This instance of Modifier claims lock on the whole class, locking all static variables
        //of the class, because alter() is static and synchronized
        SharedMemClass.alter();
    }
}
```

Volatile Variables

- **Atomic** actions are actions that occur all at once and have no chance of being interrupted or partially complete.
- This means it is impossible for **threads** to interfere with each other when performing **atomic** actions.
- In Java read and write operations on a single variable can be made **atomic** by using the access modifier **volatile**.
- **volatile** variables act as though it is enclosed in a **synchronized** block, **synchronized** on itself.
- **volatile** variables are often used to signal between **threads**.
- All operations on the variable, happen straight to the variable, it is never cached locally to a thread.

Characteristic	Synchronized	Volatile
Type of variable	Object	Object or primitive
Null allowed?	No	Yes
Can block?	Yes	No
All cached variables synchronized on access?	Yes	From Java 5 onwards
When synchronization happens	When you explicitly enter/exit a synchronized block	Whenever a volatile variable is accessed.
Can be used to combined several operations into an atomic operation?	Yes	Pre-Java 5, no. Atomic get-set of volatiles possible in Java 5.

Monitor Locks vs Semaphores

- When a **thread** enters a **synchronized** statement, it owns the **monitor lock** for the object.
- **monitor locks** restrict access to an object to one single thread.
semaphores can restrict access to an object to more than one thread.

Wait Notify

- Sometimes want a **thread** to **wait** for something in another object to happen before continuing.
- A **thread** can wait for an object to **notify** when something occurs.
- This can only be the case if the **thread** owns a **monitor lock** on the object it is waiting to be notified by and therefore must occur in a **synchronized** block.

```
if(something)//Check condition before waiting.  
    someObject.wait();  
//Continues from here after notify  
doSomethingElse();
```

- This means that the calling object will wait for someObject to call **notify** before continuing.

- The calling thread waits for an event to occur in the object called upon. The event could be from another thread or some other operating system event.
- Throws an `InterruptedException` if interrupted and not notified.
- Can be timed passing a long, similar to `sleep`.
- Must be in a synchronized block.

notify() and notifyAll()

- notify()
 - Notifies one random thread. No control in which one.
- notifyAll()
 - Notifies all waiting threads. No reason can be given for the notification.
- Both must be in synchronized blocks.

Wait/Notify-Textual Example

Producer

Creates goods at random intervals and **notifies** consumers.

- 1 Sleep for random interval.
- 2 Make a product.
- 3 **Notify** any thread waiting.

Consumer

Waits for a producer to make a good, then purchase it.

- 1 **Synchronize** on Producer.
- 2 **Wait** for a notification.
- 3 Buy Product.

Wait/Notify-Code example - Producer

```
public class Producer extends Thread {
    Product good=null;
    int nosGoods = 0;
    public void run() {
        while(!endOfSimulation()) {
//1) Sleep for random interval.
            try{
                sleep((long)(Math.random()*2000));
            }catch(InterruptedException e) {//Should not happen}
            makeGood();
        }
    }
    //Synchronized so only one good can be made at a time.
    public synchronized void makeGood() {
        if(good == null) {
//2) Make a product.
            good = new Product(nosGoods);
            nosGoods++;
            System.out.println("Making a good");
//3) Notify waiting thread.
            notify();
        }
    }
}
```

Wait/Notify-Code example - Consumer

```
public class Consumer extends Thread {
    Producer p;
    ArrayList<Product> goods;
    public Consumer(Producer p) {
        this.p = p;
        goods = new ArrayList<>();
    }
    public void run() {
        while(!p.endSimulation()) {
            if(p.hasGood()) {
//1) Synchronize Producer
                synchronize(p){
                    try{
//2) Wait for notificaation
                        p.wait();
                    }catch(InterruptedException e){//Shouldn't
                        happen}
                }
            }
//3) Buy Product.
            goods.add(p.buy());
        }
    }
}
```

Wait/Notify Common Problems

- You need to check the condition before waiting otherwise risk of never being notified.
- Waking from wait doesn't mean the condition waiting for has been satisfied, may need a loop of waiting to condition has definitely been satisfied.
- The timed wait, waits forever if given 0 as a time.
- Calling wait releases lock on object being waited on, but not on other objects locked by the class:

```
synchronized (object1);  
synchronized (object2) {  
    object2.wait(); //object2 lock released. object1 still  
                    locked  
}
```

Wait vs Sleep

- `Thread.sleep()` is a static method that sends **thread** into the "not runnable" state for a set time.
- `lock.wait()` is called on an **Object** (lock), **not** a **thread**.
 - This causes the current **thread** to wait for the **object**. **It does not cause lock to wait.**
- Before `lock.wait()` is called, the thread must **synchronize** on the lock object.
- `lock.wait()` releases the lock on the object and adds the **thread** to the "wait list" associated with lock.
- Another **thread** can **synchronize** the same lock object and call `lock.notify()` that wakes up the original, waiting **thread**.
- Essentially `wait()/notify()` is like `sleep()/interrupt()`, only the active **thread** does not need a direct reference to the sleeping thread, but only to the shared lock object.

Deadlock

- **Deadlock** describes a situation where two threads are waiting for each other to finish and never complete.
- **synchronization** is vital but it can lead to deadlock.
 - 1 A starts executing and is locked.
 - 2 B starts executing and is locked.
 - 3 A calls B and waits for B to finish
 - 4 B calls A and waits for A to finish

Starvation

- **Starvation** is where a **thread** is continually denied access to resources and is therefore unable to progress.
- This happens when "greedy" threads take resources for long periods.
- Essentially a DDOS on a **thread** and can drastically slow down a program.
- Can be caused by **threads** being given a high priority. A deterministic priority queue can cause **starvation**.

- **Threads** often act in response to an action of another thread.
- If the other **threads** action is also a response to the action of another thread, then **livelock** may occur.
- Unable to make progress.
- In this case the **threads** are not blocked, just responding to each other in a loop.
- Comparable to the awkward two people passing in a corridor going the same way each time.

Summary

- Made threaded by extends `Thread` or implements `Runnable`.
- run methods run concurrently.
- Threads claim **monitor lock** through keyword `synchronized`.
- Threads can effect each others run time behaviour through the `wait/notify` && `sleep/interrupt`. `wait/notify` requires `synchronization`.
- Presents new potential problems

The End